

# ADVERSARIAL SEARCH


Abdus Salam Azad

# Objectives

Learning how to act when the **other agents** are acting against us



# Games



# Games vs. Search

- Search – **no adversary**
  - Examples: path planning, scheduling activities
- Games – **adversary**
  - **Unpredictable opponent(s)**
  - Solution is strategy (strategy specifies move for every possible opponent reply).
  - **Time limits force an *approximate* solution**
  - **Inefficiency is intolerable**

# Zero Sum Game

- Total pay off to all players is the same for every instance of games
- Chess/Tic-tac-toe:
  - Win  $\rightarrow 1$
  - Lose  $\rightarrow -1$
  - Draw  $\rightarrow 0$

# Assumptions

- **Two agents acting alternately**
- **Utility values for each agent are the opposite of the other**
- Deterministic
- Fully observable
- Can generalize to stochastic games, multiple players, non zero-sum, etc
- In game theory terms:
  - “Deterministic, turn-taking, zero-sum games of perfect information”

# Games as search


- Initial state: e.g. board configuration of chess
- Player: which player to give the current move
- Successor function: list of (move,state) pairs specifying legal moves.
- Terminal test: Is the game finished?
- Utility function: Gives numerical value of terminal states. E.g. win (+1), lose (-1) and draw (0) in tic-tac-toe or chess

# Game Setup

- Both player wants to maximize the utility value earned at the end of the game
- We “see” the utility score as seen by MAX
- So, we can designate two players as MAX and MIN
- MAX tries to maximize its utility
- MIN tries to minimize MAX’s utility
- MAX uses search tree to determine next move.
- We play as MAX



# Partial Game Tree for Tic-Tac-Toe



# Size of search trees

- $b$  = branching factor
- $d$  = number of moves by both players
- Search tree is  $O(b^d)$
- Chess
  - $b \sim 35$
  - $D \sim 100$ 
    - search tree is  $\sim 10^{154}$  (!!)
    - completely impractical to search this
- Game-playing emphasizes being able to make optimal decisions in a finite amount of time

# Optimal Strategy

An optimal strategy leads to outcomes **at least as good** as any other strategy when the opponent plays optimally


# Strategy 1

The minimax algorithm


# The minimax algorithm

- Find the optimal *strategy* for MAX assuming an optimal MIN opponent
- Assumption: Both players play optimally!


# Two-Ply Game Tree



# Two-Ply Game Tree




# Two-Ply Game Tree





# Two-Ply Game Tree

Minimax maximizes the utility for the worst-case outcome for max



# The minimax algorithm

- Minimax value is the utility of MAX for being in the corresponding state

**MINIMAX-VALUE( $n$ ) =**

**UTILITY( $n$ )**

**If  $n$  is a terminal**

**$\max_{s \in \text{successors}(n)}$**

**MINIMAX-VALUE( $s$ )**

**If  $n$  is a max node**

**$\min_{s \in \text{successors}(n)}$**

**MINIMAX-VALUE( $s$ )**

**If  $n$  is a min node**

# Minimax algorithm

**function** MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

**return** the *action* in SUCCESSORS(*state*) with value *v*

---

**function** MAX-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

**return** *v*

---

**function** MIN-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)


$v \leftarrow \infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

**return** *v*

# Minimax is done depth-first



# Properties of Minimax

- Complete? Yes (if tree is finite)
- Optimal? Yes (against an optimal opponent)
- Time complexity?  $O(b^m)$
- Space complexity?  $O(bm)$  (depth-first exploration)
- For chess,  $b \approx 35$ ,  $m \approx 100$   
→ exact solution completely infeasible

Need to speed it up.

# Strategy 2

## Alpha Beta Pruning

# Strategy 2

## Alpha Beta Pruning

**We need not visit every node**

# Alpha-Beta Procedure

- The alpha-beta procedure can speed up a depth-first minimax search.
- **Alpha:** a lower bound on the value that a max node may ultimately be assigned

$$v \geq \alpha$$

- **Beta:** an upper bound on the value that a minimizing node may ultimately be assigned


$$v \leq \beta$$




# Alpha-Beta Procedure

- At MAX node, compare the utility estimate of current node ( $v$ ), with the  $\beta$  value, and reason whether MIN will let the game to follow this path, if not prune (i.e.  $v \geq \beta$ )
- At MIN node, compare the utility estimate of current node ( $v$ ), with the  $\alpha$  value, and reason whether MAX will let the game to follow this path, if not prune (i.e.  $v \leq \alpha$ )


# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example



# $\alpha$ - $\beta$ pruning example



# $\alpha$ - $\beta$ pruning example



# Properties of $\alpha$ - $\beta$

- Pruning **does not** affect final result. This means that it **gets the exact same result as does full minimax**.
- Good move ordering improves effectiveness of pruning
- With "perfect ordering," time complexity =  $O(b^{m/2})$   
→ **doubles** depth of search



# The $\alpha$ - $\beta$ algorithm

**function** ALPHA-BETA-SEARCH(*state*) *returns an action*

**inputs:** *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

**return** the *action* in SUCCESSORS(*state*) with value  $v$

---

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) *returns a utility value*

**inputs:** *state*, current state in game

$\alpha$ , the value of the best alternative for MAX along the path to *state*

$\beta$ , the value of the best alternative for MIN along the path to *state*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for**  $a, s$  in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

**if**  $v \geq \beta$  **then return**  $v$

$\alpha \leftarrow \text{MAX}(\alpha, v)$


**return**  $v$

# The $\alpha$ - $\beta$ algorithm


```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to state
            $\beta$ , the value of the best alternative for MIN along the path to state

  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```


# Example



# Example



# Example



**What if the depth is still very large???**



What if the depth is  
large

**Limit** the search only to  
a certain depth and use  
a **heuristic** to provide an  
estimate of the  
**MINIMAX value** of  
that node



# Some heuristics for Tic-Tac-Toe

- +100 for EACH 3-in-a-line for computer.
- +10 for EACH two-in-a-line (with a empty cell) for computer.
- +1 for EACH one-in-a-line (with two empty cells) for computer.
- Negative scores for opponent, i.e., -100, -10, -1 for EACH opponent's 3-in-a-line, 2-in-a-line and 1-in-a-line.
- 0 otherwise (empty lines or lines with both computer's and opponent's seeds).



# Resource

- Chapter 5
  - 5.1, 5.2, 5.3, 5.4
- Animation
  - <http://alphabetalekskamko.com/>
  - <http://homepage.ufp.pt/jtorres/ensino/ia/alfabeta.html>